

CS260: Machine Learning Algorithms

Lecture 1: Overview

Cho-Jui Hsieh
UCLA

Jan 7, 2019

Course Information

- Website: http://web.cs.ucla.edu/~chohsieh/teaching/CS260_Winter2019/main.html
- My office: EVI 284
- Office hours: Wednesday 11am–noon
- Online office hour: TBD
- TA: Patrick Chen (patrickchen@g.ucla.edu)
- TA for online course: Minhao Cheng (mhcheng@ucla.edu)

Course Information

- There is no textbook. Most of the topics are covered in “Deep Learning” (by Goodfellow, Bengio, Courville)

Course Information

- There is no textbook. Most of the topics are covered in “Deep Learning” (by Goodfellow, Bengio, Courville)
- Part I (basic concepts):
 - Linear models (regression, classification, clustering, dimension reduction)
 - Basic learning theory (overfitting, regularization)
- Part II (Nonlinear models):
 - Kernel methods
 - Tree-based methods
 - Deep networks
 - Applications in computer vision and NLP

Grading Policy

- Midterm exam (30%)
- Homework (30%)
 - 3 homeworks
- Final project (40%)

Final project

- Group of ≤ 4 students.
- Work on some research projects:
 - Solve an interesting problem
 - Develop a new algorithm
 - Compare state-of-the-art algorithms on some problems
 - ...
- I'll recommend some topics in the course. Feel free to discuss with me in advance.

Machine Learning: Overview

From learning to machine learning

- What is learning?

observations \rightarrow Learning \rightarrow Skill

- Skill: how to make decision (action)
 - Classify an image
 - Translate a sentence from one language to another
 - ...

From learning to machine learning

- What is learning?

observations \rightarrow Learning \rightarrow Skill

- Skill: how to make decision (action)
 - Classify an image
 - Translate a sentence from one language to another
 - ...
- Machine learning:

data \rightarrow Machine Learning \rightarrow Skill (decision rules)

Automatic the learning process!

Credit Approval Problem

- Customer record (features):

age	23 years
gender	female
annual salary	NTD 1,000,000
year in residence	1 year
year in job	0.5 year
current debt	200,000

- To be learned:

“Should we approve the credit card application? ”

Credit Approval Problem

- Customer record (features):

age	23 years
gender	female
annual salary	NTD 1,000,000
year in residence	1 year
year in job	0.5 year
current debt	200,000

- To be learned:

“Should we approve the credit card application? ”

- Data: A collection of feature-label pairs:

(customer1 feature, Yes), (customer2 feature, No), ...

Credit Approval Problem

- Customer record (features):

age	23 years
gender	female
annual salary	NTD 1,000,000
year in residence	1 year
year in job	0.5 year
current debt	200,000

- To be learned:

“Should we approve the credit card application? ”

- Data: A collection of feature-label pairs:


(customer1 feature, Yes), (customer2 feature, No), ...

- Learned model: Some decision rule

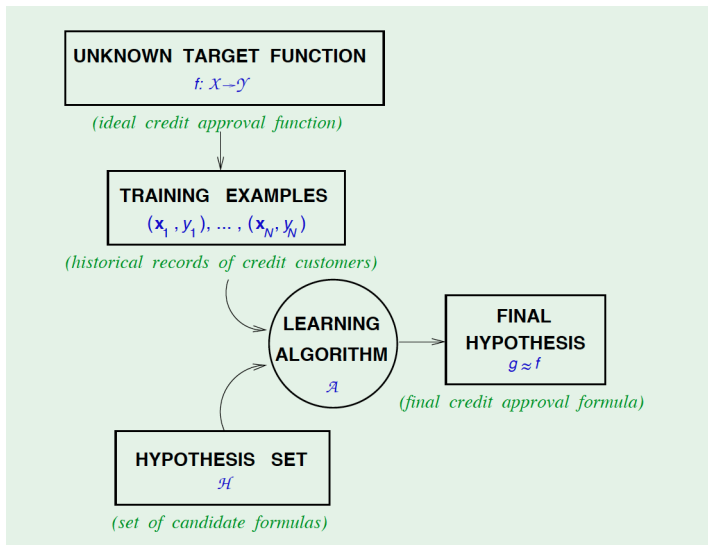
e.g., salary > 1M

Formalize the Learning Problem

- Input: $\mathbf{x} \in \mathcal{X}$ (customer application)
e.g., $\mathbf{x} = [23, 1, 1000000, 1, 0.5, 200000]$
- Output: $y \in \mathcal{Y}$ (approve/disapprove)
- Target function to be learned:
 $f : \mathcal{X} \rightarrow \mathcal{Y}$ (ideal credit approval formula)
- Data (historical records in bank):
 $\mathcal{D} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
- Hypothesis (model)
 $g : \mathcal{X} \rightarrow \mathcal{Y}$ (**learned** formula to be used)



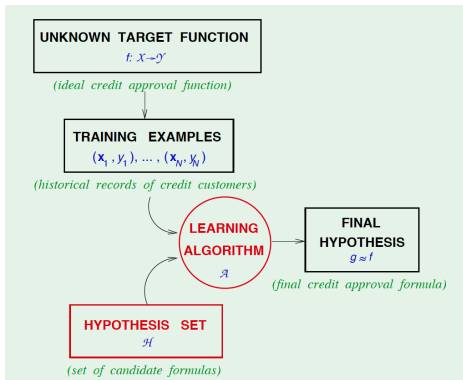
Basic Setup of Learning Problem



(Figure from "Learning from Data")

Learning Model

- A learning model has two components:
 - The **hypothesis set** \mathcal{H} :
Set of candidate hypothesis (functions)
 - The **learning algorithm**:
To pick a hypothesis (function) from the \mathcal{H}
Usually **optimization algorithm** (choose the best function to minimize the **training error**)



Perceptron

- Our first ML model: perceptron (1957)
 - Learning a linear function
 - Single layer neural network
- Next, we introduce two components of perceptron:
 - What's the hypothesis space?
 - What's the learning algorithm?

Perceptron Hypothesis Space

Define the hypothesis set \mathcal{H}

- For input $x = (x_1, \dots, x_d)$ “attributes of a customer”

Approve credit if $\sum_{i=1}^d w_i x_i > \text{threshold},$

Deny credit if $\sum_{i=1}^d w_i x_i < \text{threshold}$

- Define $\mathcal{Y} = \{+1(\text{good}), -1(\text{bad})\}$
- Linear hypothesis space \mathcal{H} : all the h with the following form

$$h(\mathbf{x}) = \text{sign}\left(\sum_{i=1}^d w_i x_i - \text{threshold}\right)$$


(perceptron hypothesis)

Perceptron Hypothesis Space (cont'd)

- Introduce an artificial coordinate $x_0 = -1$ and set $w_0 = \text{threshold}$

$$h(\mathbf{x}) = \text{sign}\left(\sum_{i=1}^d w_i x_i - \text{threshold}\right) = \text{sign}\left(\sum_{i=0}^d w_i x_i\right) = \text{sign}(\mathbf{w}^T \mathbf{x})$$

(vector form)



- Customer features \mathbf{x} : points on \mathbb{R}^d (d dimensional space)
- Labels y : $+1$ or -1
- Hypothesis h : linear hyperplanes

Select the best one from \mathcal{H}

- \mathcal{H} : all possible linear hyperplanes
- How to select the best one?

Select the best one from \mathcal{H}

- \mathcal{H} : all possible linear hyperplanes
- How to select the best one?

Find g such that $g(\mathbf{x}_n) \approx f(\mathbf{x}_n) = y_n$ for $n = 1, \dots, N$

Select the best one from \mathcal{H}

- \mathcal{H} : all possible linear hyperplanes
- How to select the best one?

Find g such that $g(\mathbf{x}_n) \approx f(\mathbf{x}_n) = y_n$ for $n = 1, \dots, N$

- Naive approach:

Test all $h \in \mathcal{H}$ and choose the best one minimizing the “training error”

$$\text{training error} = \frac{1}{N} \sum_{n=1}^N I(h(\mathbf{x}_n) \neq y_n)$$

($I(\cdot)$: indicator)

- Difficult: \mathcal{H} is of infinite size

Perceptron Learning Algorithm

Perceptron Learning Algorithm (PLA)

Initial from some \mathbf{w} (e.g., $\mathbf{w} = \mathbf{0}$)



For $t = 1, 2, \dots$

- Find a **misclassified** point $n(t)$:


$$\text{sign}(\mathbf{w}^T \mathbf{x}_{n(t)}) \neq y_{n(t)}$$

- Update the weight vector:

$$\mathbf{w} \leftarrow \mathbf{w} + y_{n(t)} \mathbf{x}_{n(t)}$$



PLA




Iteratively


- Find a misclassified point
- Rotate the hyperplane according to the misclassified point

Perceptron Learning Algorithm


- Converge for “linearly separable” case:
 - Linearly separable: there exists a perceptron (linear) hypothesis f with 0 training error
 - PLA is guaranteed to obtain f
(Stop when no more misclassified point)



(linear separable)



(not linear separable)



(not linear separable)

Binary classification

- Data:
 - Features for each training example: $\{\mathbf{x}_n\}_{n=1}^N$, each $\mathbf{x}_n \in \mathbb{R}^d$
 - Labels for each training example: $y_n \in \{+1, -1\}$
- Goal: learn a function $f : \mathbb{R}^d \rightarrow \{+1, -1\}$
- Examples:
 - Credit **approve/disapprove**
 - Email **spam/not-spam**
 - patient **sick/not sick**
 - ...


Other types of output space - Regression

- Regression: $y_n \in \mathbb{R}$ (output is a real number)
- Example:
 - Stock price prediction
 - Movie rating prediction
 - ...

Other types of output space - Multi-class prediction

Multi-class classification:

- $y_n \in \{1, \dots, C\}$ (C -way classification)
- Example: Coin recognition
 - Classify coins by two features (size, mass) ($x_n \in \mathbb{R}^2$)
 - $y_n \in \mathcal{Y} = \{1c, 5c, 10c, 25c\}$
($\mathcal{Y} = \{1, 2, 3, 4\}$)
- Other examples: hand-written digits, ...



Other types of output space - Multi-label prediction

- Multi-class problem: Each sample only has **one label**
- Multi-label problem: Each sample can have **multiple labels**

Other types of output space - Multi-label prediction

- Multi-class problem: Each sample only has **one label**
- Multi-label problem: Each sample can have **multiple labels**
- Example:
 - Document categorization (news/sports/economy/...)
 - Document/image tagging
 - ...

Other types of output space - Multi-label prediction

- Multi-class problem: Each sample only has **one label**
- Multi-label problem: Each sample can have **multiple labels**
- Example:
 - Document categorization (news/sports/economy/...)
 - Document/image tagging
 - ...
- **Extreme classification** (large output space problems):
 - Millions of billions of labels (but usually each sample only has few labels)
 - Recommendation systems: Predict a subset of preferred items for each user
 - Document retrieval or search: Predict a subset of related articles for a query

Other types of output space - structure predict

- Output as exponential

I love ML
└──┬──┘ └──┬──┘ └──┬──┘
pronoun *verb* *noun*

- Multiclass classification for each word (word \Rightarrow word class)
(not using information of the whole sentence)
- Structure prediction problem:
sentence \Rightarrow structure (class of each word)
- Other examples: speech recognition, image captioning, machine translation, ...



1. A red stop sign sitting on the side of a road.
2. A stop sign on the corner of a street.
3. A red stop sign sitting on the side of a street.

Machine Learning Problems


Machine learning problems can usually be categorized into

- Supervised learning: every \mathbf{x}_n comes with y_n (label)
(semi-supervised learning)
- Unsupervised learning: only \mathbf{x}_n , no y_n
- Reinforcement learning:


Examples contain (input, some output, grade for this output)

Unsupervised Learning (no y_n)

- Clustering: given examples x_1, \dots, x_N , classify them into K classes
- Other unsupervised learning:
 - Outlier detection: $\{x_n\} \Rightarrow \text{unusual}(x)$
 - Dimensional reduction
 - ...




supervised multiclass classification




unsupervised multiclass classification
 \longleftrightarrow 'clustering'

Semi-supervised learning


- Only some (few) x_n has y_n
- Labeled data is much more expensive than unlabeled data



supervised




semi-supervised



unsupervised (clustering)

Reinforcement Learning

- Used a lot in game AI, robotic controls
 - Agent observe state S_t
 - Agent conduct action A_t
(ML model, based on input S_t)
 - Environment gives agent reward R_t
 - Environment gives agent next state S_{t+1}
- Only observe “grade” for a certain action (best action is not revealed)
- Ads system: (customer, ad choice, click or not)



Conclusions

- Basic concept of learning:
 - Set up a hypothesis space (potential functions)
 - Define an error measurement (define the quality of each function based on data)
 - Develop an algorithm to choose a good hypothesis based on the error measurement (optimization)
- A perceptron algorithm (linear classification)
- Binary classification, multiclass, multilabel, structural prediction
- Supervised vs unsupervised learning

Questions?