

# CS260: Machine Learning Algorithms

## Lecture 10: Neural Networks


Cho-Jui Hsieh  
UCLA

Feb 20, 2019

# Neural Networks


## Another way to introduce nonlinearity

- How to generate this nonlinear hypothesis?




# Another way to introduce nonlinearity

- How to generate this nonlinear hypothesis?



- Combining multiple linear hyperplanes to construct nonlinear hypothesis

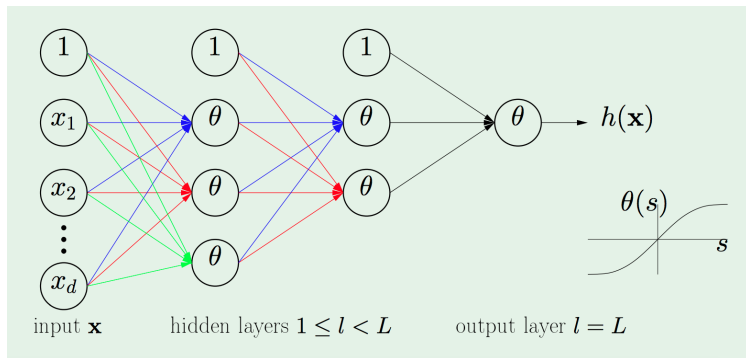


# Neural Network

- Input layer:  $d$  neurons (input features)
- Neurons from layer 1 to  $L$ : Linear combination of previous layers + activation function

$$\theta(\mathbf{w}^T \mathbf{x}), \quad \theta : \text{activation function}$$


- Final layer: one neuron  $\Rightarrow$  prediction by  $\text{sign}(h(\mathbf{x}))$



# Activation Function


## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




## tanh

$$\tanh(x)$$



## ReLU

$$\max(0, x)$$



# Formal Definitions

$$w_{ij}^{(l)} \quad \begin{cases} 1 \leq l \leq L & : \text{layers} \\ 0 \leq i \leq d^{(l-1)} & : \text{inputs} \\ 1 \leq j \leq d^{(l)} & : \text{outputs} \end{cases}$$

# Formal Definitions

$$w_{ij}^{(l)} \quad \begin{cases} 1 \leq l \leq L & : \text{layers} \\ 0 \leq i \leq d^{(l-1)} & : \text{inputs} \\ 1 \leq j \leq d^{(l)} & : \text{outputs} \end{cases}$$

$j$ -th neuron in the  $l$ -th layer:

$$x_j^{(l)} = \theta(s_j^{(l)}) = \theta\left(\sum_{i=0}^{d^{(l-1)}} w_{ij}^{(l)} x_i^{(l-1)}\right)$$



# Formal Definitions

$$w_{ij}^{(l)} \quad \begin{cases} 1 \leq l \leq L & : \text{layers} \\ 0 \leq i \leq d^{(l-1)} & : \text{inputs} \\ 1 \leq j \leq d^{(l)} & : \text{outputs} \end{cases}$$


$j$ -th neuron in the  $l$ -th layer:

$$x_j^{(l)} = \theta(s_j^{(l)}) = \theta\left(\sum_{i=0}^{d^{(l-1)}} w_{ij}^{(l)} x_i^{(l-1)}\right)$$


Output:

$$h(\mathbf{x}) = x_1^{(L)}$$

# Forward propagation




# Forward propagation



features for one data point


$$\mathbf{x} = [x_1, x_2, x_3]$$

# Forward propagation




$$x_1^{(1)} = \theta(\sum_{i=1}^3 w_{i1}^{(1)} x_i^{(0)})$$

# Forward propagation




$$x_2^{(1)} = \theta\left(\sum_{i=1}^3 w_{i2}^{(1)} x_i^{(0)}\right)$$

# Forward propagation




# Forward propagation



$$\mathbf{x}^{(1)} = \begin{bmatrix} x_1^{(1)} \\ x_2^{(1)} \\ x_3^{(1)} \end{bmatrix} = \theta \left( \begin{bmatrix} w_{11}^{(1)} & w_{21}^{(1)} & w_{31}^{(1)} \\ w_{12}^{(1)} & w_{22}^{(1)} & w_{32}^{(1)} \\ w_{13}^{(1)} & w_{23}^{(1)} & w_{33}^{(1)} \end{bmatrix} \times \begin{bmatrix} x_1^{(0)} \\ x_2^{(0)} \\ x_3^{(0)} \end{bmatrix} \right) = \theta(\mathbf{W}_1 \mathbf{x}^{(0)})$$


# Forward propagation



$$x_1^{(2)} = \theta(\sum_{i=1}^3 w_{i1}^{(2)} x_i^{(1)})$$




# Forward propagation




$$x_2^{(2)} = \theta(\sum_{i=1}^3 w_{i2}^{(2)} x_i^{(1)})$$

# Forward propagation



$$\mathbf{x}^{(2)} = \begin{bmatrix} x_1^{(2)} \\ x_2^{(2)} \\ x_3^{(2)} \end{bmatrix} = \theta \left( \begin{bmatrix} w_{11}^{(2)} & w_{21}^{(2)} & w_{31}^{(2)} \\ w_{12}^{(2)} & w_{22}^{(2)} & w_{32}^{(2)} \\ w_{13}^{(2)} & w_{23}^{(2)} & w_{33}^{(2)} \end{bmatrix} \times \begin{bmatrix} x_1^{(1)} \\ x_2^{(1)} \\ x_3^{(1)} \end{bmatrix} \right) = \theta(W_2 \mathbf{x}^{(1)})$$

# Forward propagation



$$\begin{aligned} h(\mathbf{x}) &= x_1^{(4)} = \theta(W_4 \mathbf{x}^{(3)}) = \theta(W_4 \theta(W_3 \mathbf{x}^{(2)})) \\ &= \dots = \theta(W_4 \theta(W_3 \theta(W_2 \theta(W_1 \mathbf{x})))) \end{aligned}$$

# Stochastic Gradient Descent

- All the weights  $W = \{W_1, \dots, W_L\}$  determine  $h(\mathbf{x})$

# Stochastic Gradient Descent

- All the weights  $W = \{W_1, \dots, W_L\}$  determine  $h(\mathbf{x})$
- Loss on example  $(\mathbf{x}_n, y_n)$  is

$$e(h(\mathbf{x}_n), y_n) = e(W)$$

# Stochastic Gradient Descent

- All the weights  $W = \{W_1, \dots, W_L\}$  determine  $h(\mathbf{x})$
- Loss on example  $(\mathbf{x}_n, y_n)$  is

$$e(h(\mathbf{x}_n), y_n) = e(W)$$


- To implement SGD, we need the gradient

$$\nabla e(W) : \left\{ \frac{\partial e(W)}{\partial w_{ij}^{(l)}} \right\} \text{ for all } i, j, l$$

# Computing Gradient $\frac{\partial e(W)}{\partial w_{ij}^{(l)}}$

- Use **chain rule**:

$$\frac{\partial e(W)}{\partial w_{ij}^{(l)}} = \frac{\partial e(W)}{\partial s_j^{(l)}} \times \frac{\partial s_j^{(l)}}{\partial w_{ij}^{(l)}}$$



$$s_j^{(l)} = \sum_{i=1}^d x_i^{(l-1)} w_{ij}^{(l)}$$

- We have  $\frac{\partial s_j^{(l)}}{\partial w_{ij}^{(l)}} = x_i^{(l-1)}$


# Computing Gradient $\frac{\partial e(W)}{\partial w_{ij}^{(l)}}$

- Define  $\delta_j^{(l)} := \frac{\partial e(W)}{\partial s_j^{(l)}}$
- Compute by **layer-by-layer**:

$$\begin{aligned}\delta_i^{(l-1)} &= \frac{\partial e(W)}{\partial s_i^{(l-1)}} \\&= \sum_{j=1}^d \frac{\partial e(W)}{\partial s_j^{(l)}} \times \frac{\partial s_j^{(l)}}{\partial x_i^{(l-1)}} \times \frac{\partial x_i^{(l-1)}}{\partial s_i^{(l-1)}} \\&= \sum_{j=1}^d \delta_j^{(l)} \times w_{ij}^{(l)} \times \theta'(s_i^{(l-1)}),\end{aligned}$$

where  $\theta'(s) = 1 - \theta^2(s)$  for  $\tanh$

- $\delta_i^{(l-1)} = (1 - (x_i^{(l-1)})^2) \sum_{j=1}^d w_{ij}^{(l)} \delta_j^{(l)}$






# Final layer

(Assume square loss)

- $e(W) = (x_1^{(L)} - y_n)^2$   
 $x_1^{(L)} = \theta(s_1^{(L)})$
- So,


$$\begin{aligned}\delta_1^{(L)} &= \frac{\partial e(W)}{\partial s_1^{(L)}} \\ &= \frac{\partial e(W)}{\partial x_1^{(L)}} \times \frac{\partial x_1^{(L)}}{\partial s_1^{(L)}} \\ &= 2(x_1^{(L)} - y_n) \times \theta'(s_1^{(L)})\end{aligned}$$

# Backward propagation




$$\delta_1^{(4)} = 2(x_1^{(4)} - y_n) \times (1 - (x_1^{(4)})^2)$$

# Backward propagation




$$\delta_1^{(3)} = (1 - (x_1^{(3)})^2) \times \delta_1^{(4)} \times w_{11}^{(4)}$$

# Backward propagation




$$\delta_2^{(3)} = (1 - (x_2^{(3)})^2) \times \delta_1^{(4)} \times w_{21}^{(4)}$$

# Backward propagation




$$\delta_3^{(3)} = (1 - (x_3^{(3)})^2) \times \delta_1^{(4)} \times w_{31}^{(4)}$$

# Backward propagation




$$\delta_1^{(2)} = (1 - (x_1^{(2)})^2) \sum_{j=1}^3 \delta_j^{(3)} w_{1j}^{(3)}$$

# Backward propagation



$$\delta_2^{(2)} = (1 - (x_2^{(2)})^2) \sum_{j=1}^3 \delta_j^{(3)} w_{2j}^{(3)}$$

# Backward propagation






# Backpropagation

## SGD for neural networks

- Initialize all weights  $w_{ij}^{(l)}$  **at random**
- For  $\text{iter} = 0, 1, 2, \dots$ 
  - **Forward**: Compute all  $x_j^{(l)}$  from input to output
  - **Backward**: Compute all  $\delta_j^{(l)}$  from output to input
  - Update all the weights  $w_{ij}^l \leftarrow w_{ij}^{(l)} - \eta x_i^{(l-1)} \delta_j^{(l)}$




# Backpropagation

- Just an automatic way to apply **chain rule** to compute gradient
- Auto-differentiation (AD) — as long as we define derivative for each **basic function**, we can use AD to compute any of their compositions
- Implemented in most deep learning packages  
(e.g., pytorch, tensorflow)

# Multiclass Classification

- $K$  classes:  $K$  neurons in the final layer
- Output of each  $f_i$  is the score of class  $i$

Taking  $\arg \max_i f_i(x)$  as the prediction



# Multiclass loss

- Softmax function: transform output to probability:

$$[f_1, \dots, f_K] \rightarrow [p_1, \dots, p_K]$$

where  $p_i = \frac{e^{f_i}}{\sum_{j=1}^K e^{f_j}}$

- Cross-entropy loss:

$$L = - \sum_{i=1}^K y_i \log(p_i)$$

where  $y_i$  is the  $i$ -th label

# Conclusions

- Neural network
- Back-propagation for computing gradient

Questions?